

PLD LISP

Torben Mogensen

May 26, 2021

PLD LISP is a simple LISP variant designed for “Programming Language Design”. While it is similar to other statically-scoped LISP variants (such as Scheme or Clojure), it differs from these by having pattern matching as an integral part of the `lambda` construct, and by supporting both statically and dynamically scoped functions.

This document describes PLD LISP and how to compile and run it.

1 S-Expressions

The values used in PLD LISP are *S-expressions*. An S-expression can be one of the four following forms:

A number. A number is an integer with the same range as integers in F#. Numbers are read and printed as F# integers. When reading, hexadecimal notation is allowed.

A symbol. A symbol is a name constructed from letters (Danish alphabet, both upper and lower case), digits, and the following special characters: `+-*/%=<>?!#,:{};`. Any such sequence that is not recognized as an integer constant is a valid symbol. For example, the sequence `0x3f` is a number, but `0x3g` is a symbol.

(). (pronounced “nil”) is used to represent the empty list, the false boolean value, or whatever else a programmer wants it to represent (since PLD LISP is dynamically typed).

A pair. A pair of two S-expressions s_1 and s_2 is read and written as $(s_1 . s_2)$, but there are a number of abbreviated forms used when reading or printing pairs. See below.

A closure. A closure is a pair of an S-expression and an environment (a binding from symbols to S-expressions). A closure can not be read (only constructed), and it is written as `#closure#`. The only use of a closure is applying it to a list of arguments.

We will, when reading and printing S-expressions, use the usual LISP shorthands: $(s_1 s_2 \dots s_n)$ means the same as $(s_1 . (s_2 \dots (s_n . ())))$. For example, `(car x)` is short for `(car . (x . ()))`. The notation using spaces can be combined with the dot notation, so `(a b . c)` is short for `(a . (b . c))`. Additionally, we use the shorthand `'s` to mean `(quote s)`, which can also be written (or read) as `(quote . (s))` or `(quote . (s . ()))`. In general, every time you see a period followed by an opening parenthesis, the period, the opening parenthesis, and the matching closing parenthesis can be omitted (and will be when PLD LISP shows values). Example: `(2 3 . (4))` will be shown as `(2 3 4)`. Conversely, whenever you have two values x and y separated by whitespace and followed by a closing parenthesis, you can enclose y in parentheses and add a period in front of these. For example, `(... x y)` can be expanded to `(... x . (y))`.

Note that these are only shorthands in the same way that, in F#, `[1;2;3]` is a shorthand for `1::(2::(3::[]))`.

The equivalences between S-expressions in dot-notation and their shorthands can be a bit confusing at first, but you will get used to them.

An S-expression of the form $(s_1 s_2 \dots s_n)$ usually represents a list of the elements s_1, s_2, \dots, s_n (since PLD LISP is dynamically typed, it can represent something else also). Values that can not be written in this form are not considered proper lists. For example, `(1 . 2)` is not a proper list, but `(1 . (2))` is a proper list because it can be written as `(1 2)`. We consider `()` to be a proper list (the empty list).

When reading an S-expression, it can span several lines, but only whitespace is allowed on a line after a completed S-expression. So the following is valid input:

```
( 1 2
  (3 )   4)
```

that represents the list (1 2 (3) 4), but the following is not:

```
( 1 2
  (3 )   4) 5
```

because the line that ends the S-expression that represents the list (1 2 (3) 4) has non-blank characters (namely 5) after the closing bracket. Comments start with a backslash and span to the end of the line, so

```
( 1 2 \ the first two elements
  \ followed by
  (3 )   4) \ the last two elements
```

is legal input.

The module `Sexp.fs` defines an F# data structure that represents S-expressions and functions for reading and printing S-expressions.

2 PLD LISP

The syntax of PLD LISP is represented as S-expressions. This means that *some* S-expressions in *some* contexts can be seen as expressions that are evaluated to values (that are also S-expressions). A context that evaluates S-expressions is the read-eval-print loop (REPL) in PLD LISP. Since PLD LISP syntax is represented as S-expressions, the equivalences and shorthands introduced above hold also for expressions that are evaluated. For example, (+ 2 3) is equivalent to (+ . (2 . (3))) and will evaluate to 5. (+ 2 3) is a list of three elements, but when evaluated by PLD LISP, it evaluates to a number. The fact that the syntax of an PLD LISP expression is a valid value makes it easy to make programs that manipulate PLD LISP syntax, for example by checking if a value is a syntactically correct expression.

Not all S-expressions are valid expressions in PLD LISP, though. For example, (quote . 3) is not. This will, however, not be reported until evaluation of (quote . 3) is attempted. In this sense, syntax checking is similar to dynamic type checking: It is done at runtime and only for parts of the program that are reached during execution. In most cases, PLD LISP syntax uses proper lists.

We go through the valid forms and their evaluation below:

- () evaluates to itself.
- A number evaluates to itself.
- A symbol x can be a predefined operator or keyword, in which case it evaluates to itself, or it can be a variable that evaluates to whatever x is bound to in the current environment, which consists of the global environment extended with the current local environment. Static scoping is used, so the local environment takes precedence over the global environment. If a symbol is neither a predefined operator, nor a keyword, nor a variable, evaluating the symbol gives a run-time error.

Predefined operators and keywords are

quote, lambda, lambdaD, define, if, let, load, and save	are keywords
number? and symbol?	are predefined unary operators (predicates)
cons, apply, /, %, and !=	are binary operators
+, -, *, <, =, <=, >, and >=	are variadic operators, meaning they can be applied to any number of arguments

- How an expression of the form $(e_1\ e_2\ \dots\ e_n)$ is interpreted depends on the (unevaluated) form of e_1 as follows:

- (`quote s`) evaluates to s . Note that `'s` is a shorthand for (`quote s`), so it too evaluates to s .
- (`define x e`) evaluates e to s , binds x to s in the global environment, and returns x . If x is already bound, the new binding overwrites the old.
- (`cons e1 e2`) evaluates e_1 to a and e_2 to d and returns the value $(a . d)$.
- (`load f`) loads a sequence of expressions from the file `f.le`, evaluates these, and then returns `()`. The effect is the same as if the expressions are typed in from the REPL (see below), except that `()` is returned.
- (`save f`) saves the current global environment in the file `f.le`. The saved environment is represented as a list of definitions (using `define`), one per line, so they can be read back in using `load`.
- (`if e1 e2 ... en`) is a multi-way conditional. The behaviour depends on the number of arguments n to `if`. If $n = 0$, it returns `()`. If $n = 1$, it evaluates e_1 and returns its value. If $n \geq 2$, it first evaluates e_1 to a value v . If $v \neq ()$, it then evaluates e_2 and returns its value. If $v = ()$, it applies `if` to the arguments after e_2 , i.e., does (`if e3 ... en`). Note that this implies that (`if e1 e2 e3`) works like `if e1 then e2 else e3` in other languages, that (`if e1 e2`) works like `if e1 then e2 else ()`, and that (`if e1 e2 e3 e4 e5`) works like `if e1 then e2 else if e3 then e4 else e5`.
- (`lambda p1 e1 ... pn en`) is a statically scoped function. It evaluates to a closure consisting of itself and the current local environment. A closure is printed as `#closure#` because environments are not S-expressions and can, hence, not be printed. Note that the expression `#closure#` does not evaluate to a closure.
- (`lambdaD p1 e1 ... pn en`) is a dynamically scoped function. It evaluates to itself.
- (`let p e1 e2`) is evaluated by evaluating e_1 in the current local environment ρ to a value v . The pattern p is matched against v . If p matches v , this builds an environment ρ' , which is used to extend ρ . e_2 is evaluated in the extended environment, and the result of this evaluation is returned. If p does not match v , an error is reported.
- (`u e`), where u is an unary operator, evaluates e and applies u to this value. The unary operator `number?` will return `()` if applied to anything other than a number, and return the number unchanged if applied to a number. The unary operator `symbol?` will return `()` if applied to anything other than a symbol, and return the symbol unchanged if applied to a symbol.
- (`b e1 e2`), where b is a binary operator, evaluates e_1 to v_1 and e_2 to v_2 and applies b to v_1 and v_2 : The binary operator `cons` will return the pair $(v_1 . v_2)$. The binary operator `apply` will apply the value v_1 to the list of arguments given in v_2 . For example, (`apply cons '(1 2)`) will return `(1 . 2)`. `/` divides v_1 by v_2 using integer division. If v_1 and v_2 are not numbers or if $v_2 = 0$, `()` is returned. `%` divides v_1 by v_2 using integer division and returns the remainder (using the semantics of the similar F# operator). If v_1 and v_2 are not numbers or if $v_2 = 0$, `()` is returned. `!=` compares the numbers v_1 and v_2 . If they are different, v_1 is returned, but if they are equal or they are not both numbers, `()` is returned.
- (`v e1 ... en`), where v is a variadic operator, evaluates each e_i to v_i and applies v to $v_1 \dots v_n$. The operator `+` adds its arguments. If applied to no arguments, it returns 0. If any arguments are not numbers, `()` is returned. The operator `-`, when applied to no arguments, returns 0. If applied to a single number n , it returns $-n$, if applied to $n_1 n_2 \dots n_m$, it returns $n_1 - n_2 - \dots - n_m$. If any argument is not a number, it returns `()`. The operator `*` multiplies its arguments. If applied to no arguments, it returns 1. If any arguments are not numbers, `()` is returned. The operators `<`, `=`, `<=`, `>`, and `>=` checks if the arguments are all numbers and sorted according to the operator. If they are, the first element is returned, otherwise `()` is returned. For example, (`< 7 9 13`) returns 7, while (`> 7 9 13`) returns `()`. If applied to no arguments or to arguments that are not numbers, comparison operators return `()`. Note that, while `!=` is a comparison operator, it is not variadic, as there is no sensible definition of a list being sorted by `!=`. So `!=` is a binary operator.
- If e_1 is none of the above, it is evaluated to a value v_1 . If v_1 is a closure that pairs an S-expression of the form (`lambda p1 e1 ... pn en`) with an environment ρ , the remaining elements v_2, \dots, v_m of the list are evaluated to values $v_2 \dots v_m$, and then the patterns $p_1 \dots p_n$ are in sequence matched against the list $(v_2 \dots v_m)$. If p_i matches, this builds

an environment ρ' , which is used to extend the closure environment to form a new local environment in which the corresponding expression e_i is evaluated. If no pattern matches, a run-time error is reported.

Note that this is similar to the behaviour when applying a F# expression of the form `(function $p_1 \rightarrow e_1 \mid \dots \mid p_n \rightarrow e_n$)` to an argument (v_1, \dots, v_m) . In PLD LISP, we just write `lambda` instead of `function` and omit the arrows and bars.

If the first element evaluates to an S-expression of the form `(lambdaD $p_1 e_1 \dots p_n e_n$)`, pretty much the same thing happens, except that instead of using an environment stored in a closure, it uses the current environment.

How patterns are matched to values is described below.

Pattern matching tries to match a pattern to a value. If the match is successful, it yields an environment that binds the variables in the pattern to values. Matching uses the following rules:

- The pattern `()` matches the value `()` and yields the empty environment.
- A pattern that is a number constant matches values equal to this number.
- A pattern that is a symbol x (not including keywords and predefined operators) matches any value s and yields an environment that binds x to s . Keywords are not allowed as patterns, but can occur in constant patterns (see below). Predefined operators can occur in patterns, but only match themselves.
- A pattern of the form `' v` matches values equal to v . Note that v can be any S-expression, including keywords and predefined operators.
- A pattern `($p_1 . p_2$)`, where p_1 and p_2 are patterns, matches a value `($s_1 . s_2$)` if p_1 matches s_1 yielding the environment ρ_1 , p_2 matches s_2 yielding the environment ρ_2 , and if a variable x is bound in both ρ_1 and ρ_2 , x must have the same value in both environments, otherwise the matching fails. The matching returns the combined environment $\rho_1 \cup \rho_2$.
- In all other cases, the pattern does not match the value.

Note that the usual shorthands for S-expressions is used in patterns too, so the pattern `($p_1 p_2$)` is equivalent to `($p_1 . (p_2 . ())$)` and so on.

3 Examples

We can define a function `car` that takes the head of a list by the expression

```
(define car (lambda ((a . d)) a))
```

Note that the pattern `((a . d))` specifies that `car` takes a single argument, which must be a pair of two values that are bound to the variables `a` and `d`. The expression `a` that follows this pattern just returns the first component of the pair. There are no other pattern/expression pairs in the `lambda`-expression, so if `car` is applied to something that is not a pair, a run-time error will be reported.

Note that while `((a . d))` matches a single element that is a pair (because the pattern is a list of one argument which is a pair), a pattern `(a . d)` matches any non-empty list and binds `a` to the head of the list and `d` to its tail. Since `car` takes a single argument, we use the first form. If we had used the second form, i.e., defined a function `f` by

```
(define f (lambda (a . d) a))
```

`f` would just return its first argument, so `(f 2 3 4)` would return 2.

In LISP, it is common to define a function `cadr` such that `(cadr x) \equiv (car (cdr x))`, and similar for `caar`, `cdar`, and `cddr`, and in many cases also for longer sequences of `as` and `ds`. We can define `cadr` directly by the definition

```
(define cadr (lambda (x) (car (cdr x))))
```

or using nested patterns by

```
(define cadr (lambda ((a ad . dd)) ad))
```

We can define `append` by

```
(define append (lambda ((          bs) bs
                        ((a . as) bs) (cons a (append as bs))))
```

Note that this uses two rules, one for the empty list and one for the non-empty list. Recursion is possible because `append` is bound in the global environment.

We can define a function `equal?` that tests whether two values are equal by the definition

```
(define equal? (lambda (x x) 'T (x y) ()))
```

If the two arguments are equal, the first pattern matches, so the symbol `T` is returned. Otherwise, `()` is returned. Note that `equal?` defines identity on all S-expressions, where the predefined operator `=` only defines equality on numbers.

We can make a function that tests whether a list of numbers is sorted by a given variadic comparison operator by

```
(define sortedBy (lambda (operator ()) 'T
                    (operator list) (apply operator list)))
```

The first case is needed because we consider empty lists to be sorted. We can call `sortedBy` by, e.g.,

```
(sortedBy < ' (1 2 3 4))
```

which will return `1` or by

```
(sortedBy > ' (1 2 3 4))
```

which will return `()`. In general, if the argument is a sorted list of numbers, a value different from `()` is returned, so the can be used as a condition in an if-expression.

We can implement list reversal by the function definition

```
(define reverse
  (lambda (as)
    (reverse as ())
    ((() bs) bs)
    ((a . as) bs) (reverse as (cons a bs)))))
```

This exploits that a function does not have a fixed number of parameters, so it implements both the usual reverse function and a function that takes two arguments and appends the reverse of the first argument to the second argument. The latter is typically used as a helper function when implementing reverse. A sequence of calls to `reverse` can be

```

      (reverse (a b c))
~> (reverse (a b c) ()) using first rule
~> (reverse (b c) (a))  using third rule
~> (reverse (c) (b a))  using third rule
~> (reverse () (c b a)) using third rule
~> (c b a)              using second rule
```

Note that the arguments are shown as values – to make the initial call from the REPL, you should write `(reverse '(a b c))`.

4 Compiling and Running PLD LISP

The file `LISP.zip` contains an interpreter written in F# for this LISP variant. The interpreter implements a read-eval-print loop (REPL) and starts with an empty global environment. You can compile the interpreter with the command `source compile.sh` and run it with the command `mono lisp.exe` or, if .NET is installed, just by `./lisp.exe` if on Linux or MacOS or `lisp.exe` if using Powershell on Windows.

Type in an S-expression at the "> " prompt. It may span several lines. When an "Enter" is typed and a complete S-expression is found, this is evaluated and the result shown after a "= ". If an error is found during parsing or evaluation, a message is displayed after a "! ", and a backtrace is shown. The backtrace shows the recursion stack by listing (in reverse order) which functions that were applied to which values to reach the error. Note that functions are shown as lambda-expressions and not as names.

In both cases, the "> " prompt is shown again. Close a session by pressing ^D (control-D) at the prompt (^Z if using Windows).

The file `listfunctions.le` contains definitions of some simple, but useful, list functions, including `car`, `cadr`, `equal?`, `length`, `append`, and `reverse`.

An example session using the PLD LISP REPL is shown below

```
$ mono lisp.exe
PLD LISP version 2.1
> (load listfunctions)
> = car
> = cdr
> = caar
> = cadr
> = cdar
> = cddr
> = list
> = length
> = append
> = reverse
> = equal
> = ()
> (append '(a b c) '(d e f))
= (a b c d e f)
> (append '(1 . 2) ())
! No patterns matched arguments (2 ())
! when applying (lambda (() bs) bs ((a . as) bs) (cons a (append as bs)))
! to (2 ())
! when applying (lambda (() bs) bs ((a . as) bs) (cons a (append as bs)))
! to ((1 . 2) ())
>
```

Note that the backtrace when applying `append` to a non-proper list (which leads to an error) shows the body of `append` instead of its name.

Infinite recursion will cause the interpreter to run out of stack, so you get a very long F# backtrace.